

AgriHGT: Heterogeneous Graph Transformer for Multi-Crop Leaf Disease Classification

Author One, Author Two, Author Three

Department of Computer Science and Engineering, East West University, Dhaka, Bangladesh

{author1, author2, author3}@example.com

Abstract—Leaf disease classification in multi-crop field images is usually treated as independent image recognition, although crops, diseases, and disease categories are related and visually similar leaves carry complementary evidence. We propose AgriHGT, a two-stage visual-relational framework that places partially fine-tuned backbone features of each leaf image in a heterogeneous graph with crop, disease, and disease-category nodes, propagates information with a Heterogeneous Graph Transformer, fuses the result with a gated residual visual path, and predicts the disease with a leaf-disease compatibility decoder. On the 28-class Agri-Vision4 dataset (Tomato, Papaya, Zucchini, Bottle Gourd), we build a duplicate-aware, leakage-controlled split through MD5 deduplication, perceptual-hash linking, and an iterative near-duplicate audit, and prevent label shortcuts by separating message-passing and supervised training nodes. Seven convolutional, transformer, and self-supervised backbones, three message-passing depths, a feature-only logistic-regression reference, and a nine-configuration component ablation are evaluated. The best configuration, [backbone, depth], reaches [xx.xx]% test accuracy and [0.xxxx] macro-F1, compared with [xx.xx]% for the corresponding feature-only reference, and the ablation shows that [component] contributes most.

Index Terms—Plant disease classification, heterogeneous graph transformer, graph neural network, self-supervised learning, DINOv3, Agri-Vision4.

I. INTRODUCTION

Leaf diseases reduce vegetable yield and quality, and in smallholder settings such as Bangladesh, expert diagnosis is often unavailable when it is needed. Deep learning on leaf images has become the standard route to automated diagnosis [?], [?], progressing from CNNs to vision transformers and self-supervised representations [?], [?], [?], [?], [?]. These models, however, classify each image independently. In a multi-crop setting, a leaf belongs to a crop, its disease belongs to a broader category, the same category spans several crops, and visually similar leaves of the same crop tend to share a diagnosis. This structure is not used by an image-only classifier.

Two lines of work address parts of it. Graph neural networks relate visual regions or leaf images [?], [?], and agricultural knowledge graphs encode crops, diseases, pests, and symptoms [?], [?], [?]. They are rarely combined, and combining them raises three practical risks: augmented copies of one photograph leaking across partitions, label information returning to supervised nodes through the graph, and conclusions that hold for only one backbone.

We therefore propose AgriHGT, which joins image features and agricultural knowledge in one heterogeneous graph pro-

cessed by the Heterogeneous Graph Transformer (HGT) [?]. Our contributions are:

- a heterogeneous graph of 26,905 nodes and 13 relation types linking leaf images to crops, diseases, and disease categories, with gated residual fusion and a leaf-disease compatibility decoder;
- a duplicate-aware, leakage-controlled protocol for Agri-Vision4 [?], combining exact deduplication, perceptual-hash linking, an iterative near-duplicate audit, and per-class support control, plus a message/supervision partition that removes label shortcuts;
- a comparison of seven backbones, three HGT depths, and a feature-only reference, testing whether graph learning helps consistently across representations; and
- a nine-configuration ablation with paired McNemar tests, and evaluation with class-balanced, chance-corrected, and calibration metrics.

II. RELATED WORK

Image-based plant disease recognition. CNN transfer learning established high accuracy on leaf images [?], [?], with ResNet [?], [?] and EfficientNet [?] models widely used. Vision and Swin transformers [?], [?], [?] and other recent designs [?], [?], [?] capture longer-range context, and self-supervised learning is increasingly used [?], [?]. All treat images independently.

Graphs for plant disease. Graph-based models capture relationships among visual regions or leaf images, including feature-space similarity graphs over images [?], [?]. Agricultural knowledge graphs represent semantic relationships among crops, diseases, pests, symptoms, and related entities [?], [?], [?]. AgriHGT combines both kinds of information in one graph.

Heterogeneous graph learning. HGT uses node- and edge-type-dependent parameters and relation-aware attention for graphs with many node and relation types [?]. AgriHGT applies it to a graph in which leaf images, crops, diseases, and categories coexist, so that similarity messages and knowledge messages are handled by separate parameters.

III. MATERIALS AND METHODS

A. Overview

AgriHGT is a two-stage visual-relational classifier (Fig. 1). First, a pretrained visual backbone is partially fine-tuned on the training partition and used to extract one fixed feature




Fig. 1. Overview of the AgriHGT methodology.




Fig. 2. Original and augmented images across the four crops of Agri-Vision4.

vector per retained leaf image. Second, these features become `leaf_image` nodes in a heterogeneous graph with crop, disease, and disease-category nodes, on which relation-aware message passing is performed with the Heterogeneous Graph Transformer (HGT) [?]. The graph representation is fused with a residual visual representation before disease prediction.

B. Dataset

Agri-Vision4 [?] contains leaf images of Tomato, Papaya, Zucchini, and Bottle Gourd in 28 crop-specific healthy and disease classes, collected in Bangladeshi fields under natural light and validated by agricultural experts. It provides original images and augmented copies (rotation, flipping, brightness and contrast adjustment, noise); augmented images form a large share of every crop (Fig. 2). Before duplicate removal there were 28,121 files, of which 5,246 were originals (Table I). Because class frequencies are unequal, accuracy is reported with macro-F1 and other class-balanced metrics.

C. Leakage-Controlled Data Partitioning

Exact duplicates. An MD5 hash was computed for all 28,121 files, and one file per identical-hash group was kept. This removed 1,255 byte-identical duplicates, leaving 26,866 images (4,304 original, 22,562 augmented).

Duplicate-aware initial split. A random split over originals and augmentations could put an image and its transformed copy in different partitions and inflate validation or test scores, so the originals were split first with class-stratified 70/15/15 sampling (3,004/650/650). Perceptual hashes with `hash_size=16` (256 bits) were then computed, and each augmented image was compared with originals of the same crop and class; if the minimum Hamming distance was at most six, it was linked to the closest original and assigned to its partition. This linked 1,552 augmented images; the other

21,010 could not be confidently associated and were assigned to training.

Near-duplicate audit. Same-class perceptual hashes were then compared across partitions for original–original, original–augmented, and augmented–augmented pairs. Whenever a cross-partition pair had Hamming distance ≤ 6 , the validation or test image was moved to training, repeating until no violation remained; successive rounds found 41, 11, 3, and 0 violating pairs.

Per-class support control. Because the previous steps could remove the only validation or test image of a minority class, a class lacking either was given an eligible training image, after which the audit was repeated; if the reassigned image created a new near-duplicate relation, it was returned and another candidate tried. The final split has 25,122 training, 896 validation, and 848 test images, with at least one validation and one test image per class, and is shared by all backbone experiments. It is *leakage-controlled* rather than leakage-free: the audit covers same-class pairs and the chosen hash and threshold, so the supported claim is that no same-class cross-partition pair exists at Hamming distance ≤ 6 .

D. Visual Backbones

Each image’s backbone feature initializes its `leaf_image` node and defines the `similar_to` relations, so the graph changes with the backbone. Seven backbones spanning convolutional, transformer, supervised, and self-supervised families used in recent plant-disease studies [?], [?], [?] were compared (Table II), both to find the strongest one and to test whether heterogeneous graph learning behaves consistently across representations. ResNet50 is a standard residual CNN baseline whose residual connections ease deep training, used widely in plant-disease transfer learning [?], [?]. DenseNet121 connects each layer to all earlier outputs, encouraging feature reuse so that small lesions and larger texture patterns both remain available. EfficientNet-B3 scales depth, width, and resolution jointly, suiting symptoms at multiple scales, and performs strongly on leaf disease [?]. ConvNeXt-Tiny is a modern CNN with transformer-inspired design, a middle point showing whether gains need attention or also arise in a modernized CNN. Swin-Tiny is a supervised hierarchical transformer with shifted-window local attention, applied to plant disease in [?], [?], [?]. DINOv2 ViT-B/14 is self-supervised, learning from unlabeled images an embedding space that may preserve visual structure beyond supervised classes, which matters because the graph needs meaningful neighborhoods, not only a strong classifier [?], [?]. DINOv3 ViT-B/16, a newer DINO model, tests whether a newer self-supervised representation helps; the DINOv2–DINOv3 comparison isolates representation generation within one family.

E. Backbone Adaptation and Feature Extraction

Generic pretrained features may emphasize shape, background, and global appearance rather than subtle color and texture cues, so each backbone was first adapted on the training partition with a temporary 28-class head, fine-tuning

TABLE I
ORIGINAL IMAGES PER CROP-SPECIFIC CLASS BEFORE DUPLICATE REMOVAL (TOTAL 5,246)

Bottle Gourd (1,658)		Papaya (1,684)		Tomato (916)		Zucchini (988)	
Alternaria Leaf Blight	303	Bacterial Blight	183	Downy	57	Angular Leaf Spot	120
Anthrachnose	276	Carica Insect Hole	318	Healthy	288	Anthrachnose	129
Downy Mildew	286	Curl'd Yellow Spot	538	Mosaic	195	Downy Mildew	153
Early Alternaria L. Blight	179	Healthy Leaf	189	Spot	311	Dry Leaf	67
Fungal Damage Leaf	39	Mosaic Virus	119	White Spot	65	Healthy	88
Healthy	260	Pathogen Symptoms	286			Insect Damage	78
Mosaic Virus	315	Yellow Necrotic Spot Holes	51			Iron Chlorosis Damage	65
						Xanthomonas Leaf Spot	86
						Yellow Mosaic Virus	202

TABLE II
VISUAL BACKBONES AND THEIR ROLE IN AGRIHGT

Backbone	Architecture	Pretrain.	Reason for inclusion
ResNet50	Residual CNN	Sup.	Standard CNN reference
DenseNet121	Dense CNN	Sup.	Feature reuse
EfficientNet-B3	Scaled CNN	Sup.	Compound scaling
ConvNeXt-Tiny	Modern CNN	Sup.	CNN-transformer bridge
Swin-Tiny	Hierarchical ViT	Sup.	Local-window attention
DINOv2 ViT-B/14	ViT	Self-sup.	Strong SSL features
DINOv3 ViT-B/16	ViT	Self-sup.	Recent DINO generation

only the last stage: layer4 for ResNet50; denseblock4, norm5, and classifier for DenseNet121; features[7:] for EfficientNet-B3; features[7] for ConvNeXt-Tiny; the final stage and terminal norm for Swin-Tiny; and the final two blocks and terminal norms for DINOv2 and DINOv3, each with its head. Partial fine-tuning limits cost and preserves low- and mid-level pretrained features on a dataset far smaller than the pretraining data. Validation accuracy was monitored every epoch (with patience-based early stopping where used), the best checkpoint was restored, the head was replaced by an identity, and the frozen backbone ran one deterministic pass without augmentation over all 26,866 images at its native resolution and normalization. This gives feature matrices of $26,866 \times d$ with $d = 2,048$ (ResNet50), 1,024 (DenseNet121), 1,536 (EfficientNet-B3), and 768 (ConvNeXt-Tiny, Swin-Tiny, DINOv2, DINOv3). Each dimension was standardized as $x' = (x - \mu_{\text{train}}) / \sigma_{\text{train}}$ with statistics from the training partition only, applied unchanged to validation and test, separately per backbone.

F. Heterogeneous Graph Construction

Graph learning has captured relations among image regions or leaves [?], [?], and agricultural knowledge graphs relate crops, diseases, pests, and symptoms [?], [?], [?]; AgriHGT combines both (Fig. 3). One graph per backbone has 26,905 nodes: 26,866 leaf_image, 4 crop, 28 disease, and 7 disease_category nodes (Healthy, Viral, Bacterial, Fungal, Pest, Nutrient Deficiency, Other), with categories fixed from class names before training [?], [?]. Leaf nodes take the standardized backbone vector; non-image nodes have no pixels and receive learnable input embeddings.

Six forward relations are used: grown_on (leaf→crop), affects (crop→disease), is_type_of (disease→category), similar_to (leaf→similar leaf), has_disease (selected training leaf→disease), and




Fig. 3. AgriHGT heterogeneous graph: leaf_image, crop, disease, and disease_category nodes and their typed relations.

co_category (disease→disease of the same category). Adding a reverse for each and a self-loop per node gives 13 relation types. For similar_to, following feature-space similarity graphs [?], each image links to its five nearest neighbors of the same crop under

$$d_{\cos}(\mathbf{x}_i, \mathbf{x}_j) = 1 - \frac{\mathbf{x}_i^T \mathbf{x}_j}{\|\mathbf{x}_i\|_2 \|\mathbf{x}_j\|_2}, \quad (1)$$

restricting to the crop to reduce the influence of crop-level leaf shape, giving 134,330 forward edges, recomputed per backbone.

Message and supervision partition. Linking every supervised training image to its true disease would leak its label along leaf_image→disease→leaf_image. The training partition was therefore split, class-stratified 75/25 with originals and linked derivatives kept together, into 18,881 train_msg images, which keep has_disease edges but are excluded from the loss, and 6,241 train_sup images, which have no has_disease edge and alone form the loss. Validation and test nodes also lack has_disease edges, so train_sup, validation, and test nodes follow the same rule. With reverse relations and self-loops the graph has 387,301 edges.

G. AgriHGT Architecture

AgriHGT has three components: heterogeneous message passing, gated fusion of graph and visual representations, and a leaf-disease compatibility decoder (Fig. 4).

Projection and HGT. Node features of type $\tau(v)$ are mapped to 256 dimensions, $\mathbf{h}_v^{(0)} = \tanh(W_{\tau(v)} \mathbf{x}_v + \mathbf{b}_{\tau(v)})$. HGT uses type-dependent transformations and relation-aware attention [?], fitting a graph in which a leaf-similarity message means something different from a crop-disease or taxonomy




Fig. 4. AgriHGT architecture: visual feature extraction, heterogeneous graph construction, relation-aware HGT message passing, residual visual fusion, and disease prediction.

message. We use 4 heads, hidden size 256, dropout 0.2, and depths $n_{\text{layers}} \in \{0, 1, 2\}$. The zero-layer model keeps the projections, learnable non-image embeddings, residual branch, fusion, and decoder but no message passing; since the original pyHGT code still built one convolution at depth 0, it was modified to perform none.

Gated residual fusion. The standardized image feature is projected as $\mathbf{r}_i = \text{LayerNorm}(W_r \mathbf{x}_i + \mathbf{b}_r)$. With graph-updated leaf representation $\mathbf{h}_i$,

$$\mathbf{g}_i = \sigma(W_g[\mathbf{r}_i \parallel \mathbf{h}_i] + \mathbf{b}_g), \quad \mathbf{z}_i = \mathbf{g}_i \odot \mathbf{r}_i + (1 - \mathbf{g}_i) \odot \mathbf{h}_i, \quad (2)$$

where $\parallel$ is concatenation. The gate bias starts at 2.0 ($\sigma(2) \approx 0.88$), so the model first favors the visual path and learns how much graph information to admit.

Compatibility decoder. Leaf and disease features are projected to 128 dimensions, $\mathbf{u}_i = W_l \mathbf{z}_i + \mathbf{b}_l$ and $\mathbf{v}_j = W_d \mathbf{d}_j$, and scored against all 28 disease nodes:

$$s_{ij} = W_o \text{Dropout}_{0.1}[\text{ReLU}(\mathbf{u}_i + \mathbf{v}_j)] + b_o, \quad (3)$$

and the prediction is $\hat{y}_i = \arg \max_j s_{ij}$, the disease whose representation best matches the leaf.

H. Training and Transductive Setting

Training is full-batch over the whole graph, with loss only on the 6,241 `train_sup` nodes. Class weights $\tilde{w}_c = 1/n_c$ over `train_sup` counts are normalized to mean one, and

$$\mathcal{L} = -\frac{1}{|\mathcal{S}|} \sum_{i \in \mathcal{S}} w_{y_i} \log \frac{\exp(s_{i, y_i})}{\sum_{j=1}^{28} \exp(s_{ij})}. \quad (4)$$

AdamW (learning rate 10^{-3} , weight decay 10^{-4}), gradient clipping at 1.0, and cosine decay toward 10^{-6} are used. ResNet50, Swin-Tiny, and DenseNet121 runs use up to 300 epochs with patience 300 (the full budget, keeping the best validation checkpoint); ConvNeXt-Tiny, EfficientNet-B3, DINOv2, and DINOv3 use up to 400 epochs with patience 50; the minimum improvement is 10^{-4} . Depths 0, 1, and 2 share one partition per backbone, and the preferred depth and checkpoint are chosen by validation accuracy only; test results never drive selection. The setting is *transductive*: validation and test leaf nodes are in the graph and take part in message passing, but their labels are excluded from the loss and they have no `has_disease` edges. Results therefore measure

classification of unlabeled nodes already in the graph, not fully inductive prediction for images added after training.

I. Reference Models and Ablation Design

Logistic-regression reference. A multinomial logistic regression (up to 2,000 iterations) is trained per backbone on the same standardized features without any graph, crop, disease, or category nodes, HGT, fusion, or decoder, reflecting conventional feature-only recognition [?], [?]. It uses all 25,122 training labels, since no label-bearing edges require a message/supervision split, so it is a visual-feature reference rather than a matched supervision ablation. **Zero-layer reference.** HGT-0 keeps the rest of AgriHGT, so comparing it with HGT-1 and HGT-2 isolates message propagation.

Component ablation. With DINOv3 features, two HGT layers (best DINOv3 validation depth) form the full reference A0, and one component is changed at a time (Table III), keeping hidden size 256, 4 heads, dropout 0.2, the MLP decoder, the weighted loss, the split, the feature bank, the `train_msg/train_sup` partition, and the validation and test sets. For A2–A5 the graph is rebuilt without only the named nodes or relations (the within-crop neighbor graph is kept for A3–A5 and removed only in A2). A6 keeps the structure but zeroes the meta-node embeddings, A7 feeds $\mathbf{h}_i$ directly to the decoder, and A8 replaces Eq. (2) with $\mathbf{z}_i = 0.5\mathbf{r}_i + 0.5\mathbf{h}_i$. A0 and A2–A8 use up to 400 epochs, AdamW (10^{-3} , 10^{-4}), cosine decay to 10^{-6} , patience 50, minimum improvement 10^{-4} , and seed 42; A1 reuses the zero-layer DINOv3 run under the same settings. With a single seed, differences are component-level observations, which the paired McNemar tests below qualify.

J. Evaluation Metrics

For N test images with per-class counts TP_c, FP_c, FN_c, TN_c , we report accuracy, precision $P_c = TP_c / (TP_c + FP_c)$, recall $R_c = TP_c / (TP_c + FN_c)$, specificity $Sp_c = TN_c / (TN_c + FP_c)$, $F1_c = 2P_c R_c / (P_c + R_c)$, and Jaccard $J_c = TP_c / (TP_c + FP_c + FN_c)$, with macro averages $\frac{1}{28} \sum_c (\cdot)$ that weight every class equally; balanced accuracy is the macro recall. Chance-corrected agreement is measured by Cohen's

TABLE III
COMPONENT ABLATION CONFIGURATIONS (DINOv3 FEATURES)

ID	Modification relative to full AgriHGT (A0, two HGT layers)
A0	None: all relations and components retained
A1	Zero HGT layers: no GeneralConv message passing
A2	Remove <code>similar_to</code> and its reverse
A3	Remove crop nodes, <code>grown_on</code> , <code>affects</code> , and <code>reverses</code>
A4	Remove category nodes, <code>is_type_of</code> , <code>co_category</code> , and <code>reverses</code>
A5	Remove <code>has_disease</code> and reverse (no visual grounding of diseases)
A6	Replace learnable crop/disease/category embeddings with zeros
A7	Remove the residual visual bypass (graph representation only)
A8	Replace the learned gate with a fixed equal-weight average

$\kappa = (p_o - p_e)/(1 - p_e)$, with $p_e = \sum_c n_c \hat{n}_c / N^2$, and by the multiclass Matthews correlation

$$MCC = \frac{C N - \sum_c \hat{n}_c n_c}{\sqrt{(N^2 - \sum_c \hat{n}_c^2)(N^2 - \sum_c n_c^2)}}, \quad (5)$$

where C is the number of correct predictions and n_c , $\hat{n}_c$ are true and predicted class counts. From predicted probabilities we report top-3 accuracy, macro one-vs-rest ROC-AUC, negative log-likelihood (NLL), and expected calibration error $ECE = \sum_b \frac{|B_b|}{N} |\text{acc}(B_b) - \text{conf}(B_b)|$ over 15 confidence bins. Accuracy has a 95% Wilson interval

$$\left(\hat{p} + \frac{z^2}{2N} \pm z \sqrt{\frac{\hat{p}(1-\hat{p})}{N} + \frac{z^2}{4N^2}} \right) / \left(1 + \frac{z^2}{N} \right), \quad z = 1.96, \quad (6)$$

macro-F1 a 95% bootstrap interval (2,000 resamples of the test set), and each ablation is compared with A0 on the same 848 images by an exact McNemar test on the discordant pairs (n_{01}, n_{10}) , $p = 2 \sum_{k \leq \min(n_{01}, n_{10})} \binom{n_{01}+n_{10}}{k} 2^{-(n_{01}+n_{10})}$. All values are computed by a released script directly from saved predictions.

K. Experimental Setup

Experiments ran on Kaggle with one NVIDIA Tesla T4 GPU. PyTorch and Torchvision served backbone training and feature extraction; DINOv2 was loaded from its official implementation and DINOv3 ViT-B/16 via `timm`; pyHGT provided the graph and HGT layers (with the zero-layer change of Section III-G); PyTorch Geometric supported graph operations; scikit-learn provided nearest-neighbor search, logistic regression, and evaluation; Pandas and NumPy handled data; and Pillow and ImageHash handled images and perceptual hashing. The split was fixed across backbones, seed 42 was used for the main and ablation experiments, and a fixed hold-out validation set (not k -fold cross-validation) selected checkpoints and depths. All experiments are single-seed; multi-seed variability was not estimated and is a limitation.

IV. RESULTS

Backbones and depth. Table IV reports every backbone with the feature-only reference and HGT depths 0–2. The strongest configuration is [backbone, depth], with [xx.xx]% test accuracy (95% CI [[xx.xx, xx.xx]]), [0.xxxx] macro-F1, and $\kappa = [0.xxxx]$. Relative to its logistic-regression reference, message passing changes accuracy by [$\pm$ x.xx] points and macro-F1 by [$\pm$ 0.xxxx]; across backbones, HGT-1/2 improves

on HGT-0 for [k] of 7 backbones. Self-supervised DINO features [do / do not] outperform the supervised CNN and Swin features.

Component ablation. Table V shows that removing [component] causes the largest macro-F1 drop ([−0.xxxx], McNemar $p = [0.xxx]$), while [components] change results by less than [x] points and are not significant at $p < 0.05$.

Per-class behavior. Table VI shows that the lowest F1 occurs for [classes], mostly minority classes, whereas [classes] are recognized almost perfectly.

V. DISCUSSION AND LIMITATIONS

The comparison with HGT-0 is the most direct test of message passing, because only propagation differs; the logistic-regression reference instead shows how much is already in the adapted features, but it sees all 25,122 training labels while AgriHGT supervises only the 6,241 `train_sup` nodes, so it is not a matched ablation. The gate initialization ($\sigma(2) \approx 0.88$) lets the model fall back on visual evidence when graph messages are unhelpful, which is why the residual-bypass ablation (A7) is informative about how much the graph alone can carry.

Four limitations qualify the results. First, the evaluation is transductive: test images take part in message passing without their labels, so performance on images added after training remains to be measured. Second, leakage control covers same-class pairs within Hamming distance six of a 256-bit perceptual hash; near-duplicates beyond this threshold or across classes are not excluded. Third, all experiments use one seed and one fixed split, so differences smaller than the Wilson and bootstrap intervals should not be over-interpreted. Fourth, the four crops come from one region, so transfer to other crops, regions, and imaging conditions is untested.

VI. CONCLUSION

AgriHGT combines backbone features of leaf images with crop, disease, and disease-category knowledge in a heterogeneous graph, processed by relation-aware HGT message passing, gated residual fusion, and a compatibility decoder. On a leakage-controlled split of Agri-Vision4 with a label-shortcut-free supervision scheme, it was evaluated across seven backbones, three depths, and nine ablations, reaching [xx.xx]% accuracy and [0.xxxx] macro-F1 with [backbone]. Future work will evaluate inductive inference, multiple seeds and splits, cross-class near-duplicate auditing, and additional crops and regions.

TABLE IV
TEST RESULTS PER BACKBONE: FEATURE-ONLY LOGISTIC REGRESSION AND AGRIHGT WITH 0-2 HGT LAYERS ($N = 848$). [†]DEPTH SELECTED BY
VALIDATION ACCURACY. ACCURACY-TYPE VALUES IN %

[illegible]

TABLE V
COMPONENT ABLATION WITH DINOv3 FEATURES ($N = 848$). Δ RELATIVE TO A0; MACRO-F1 CI BY BOOTSTRAP; p FROM THE EXACT MCNEMAR TEST AGAINST A0

[illegible]

TABLE VI
PER-CLASS ONE-VS-REST TEST RESULTS OF THE SELECTED AGRIHGT MODEL (BG = BOTTLE GOURD, PA = PAPAYA, TO = TOMATO, ZU = ZUCCHINI; n = TEST SUPPORT)

Class	Prec.	Rec.	Spec.	F1	Jacc.	MCC	 n 	 	Class	Prec.	Rec.	Spec.	F1	Jacc.	MCC	 n
<i>Per-class rows are generated from the selected model's test predictions.</i>																